

Ashish Maurya

Location: Maharashtra, Mumbai, India

[Ashishmaurya.com](https://ashishmaurya.com) | [LinkedIn](#) | [GitHub](#) | [Leetcode](#) Email: ramashishmaurya738@gmail.com | Mobile: 9769295699

SUMMARY

I am an aspiring AI Engineer with strong fundamentals in **Python, Data Processing, and Machine Learning**. I have hands-on experience building end-to-end data pipelines, AI-powered applications (LLM/RAG), and deploying scalable solutions. I am proficient in **Python, SQL, FastAPI, and LangChain**, and passionate about solving real-world problems through data-driven engineering and intelligent systems.

SKILLS

Programming: Python

Data Engineering & AI Tools: FastAPI, LangChain, LangGraph, Redis, MLflow, Celery

Cloud DevOps: AWS (EC2, S3), Docker

Databases: PostgreSQL, MongoDB, Qdrant (Vector DB)

Machine Learning & AI: Deep Learning, NLP, LLMs/RAG, Scikit-learn, PyTorch

Tools: Git, GitHub, Jupyter Notebook, PyMuPDF

EDUCATION

University of Mumbai
Bachelor of Data Science

Maharashtra, Mumbai, India
Aug 2022 – Jun 2025

PROJECTS

NeuralDocs AI | Reactjs, FastAPI, LangChain, RAG, Qdrant, Docker, WebSockets, Redis

[\[View on GitHub\]](#)

- Built a full-stack RAG assistant (**React, FastAPI**) enabling users to upload and query complex PDFs in natural language.
- Integrated real-time **WebSocket streaming with LangChain** to deliver instant, word-by-word AI responses with zero latency.
- Designed an asynchronous ingestion pipeline **using Celery and Redis for seamless, non-blocking background processing of large files.**
- Optimized retrieval performance and cost by **leveraging Qdrant vector database**, context-aware querying, and LLM caching.

Job Salary Prediction | Python, Random Forest, XGBoost, MLflow, Docker, Render, Streamlit

[\[View on GitHub\]](#)

- Built a **Job Salary Prediction system** using **Machine Learning** models (Random Forest, XGBoost) trained on real-world job listings to predict salaries based on role, experience, location, and skills with high accuracy.
- Integrated **MLflow** for complete experiment tracking — logged model parameters, metrics, and artifacts across multiple runs to compare and select the best performing model for production deployment.
- Containerized the full ML pipeline using **Docker** and deployed on **Render** for public access, enabling real-time salary predictions through an interactive **Streamlit** web application.

Email Spam Detection | Python, NLP, TF-IDF, Naive Bayes, Logistic Regression

[\[View on GitHub\]](#)

- Built an **Email Spam Detection system** using **NLP techniques** — preprocessed email text with tokenization, stop-word removal and **TF-IDF vectorization** to classify emails as spam or not spam.
- Trained and compared **Naive Bayes and Logistic Regression** models on real email datasets, achieving **95% accuracy** with a clean **Streamlit** interface for live email classification.

CERTIFICATIONS

- [Fundamental Tools of Data Wrangling](#)
- [Complete Data Science, Machine Learning, DL, NLP](#)
- [Structured Query Language \(SQL\)](#)
- [Data Structures and Algorithms Using Python](#)